

# File server with backend authentication

Major Amit Pathania 163054001

Nivia Jatain 15305R007

Indian Institute of Technology Bombay

November 16, 2017

# Table of Contents

- 1 System Architecture
- 2 Load Generator
- 3 Phase2 Bottleneck
- 4 Optimisation
- 5 Results after optimization
- 6 Summary of results
- 7 Lessons learnt

# System Architecture

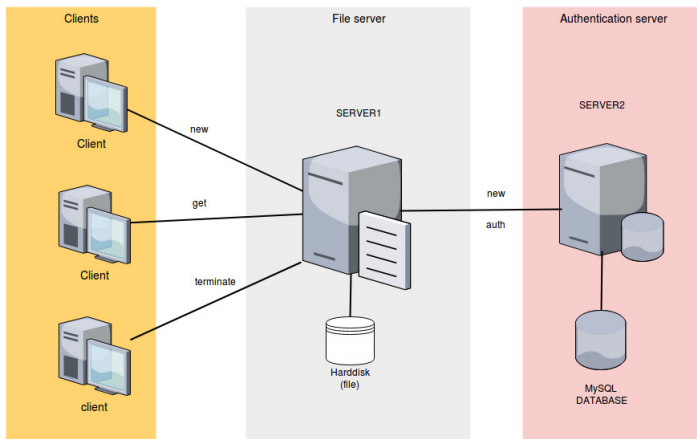
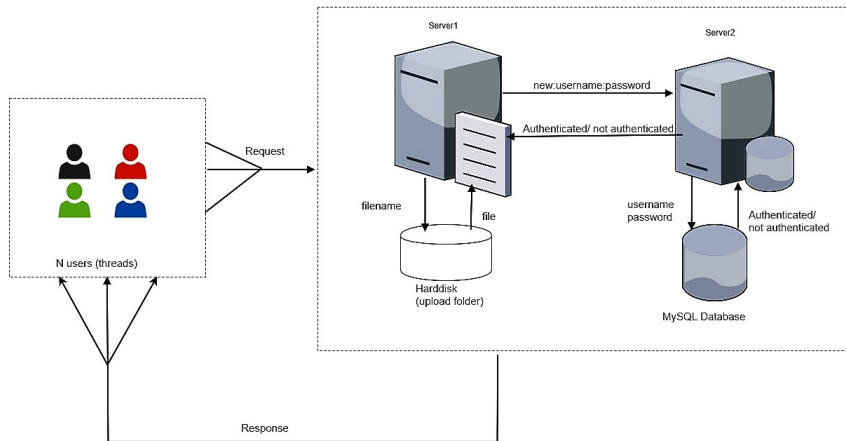


Figure: System Architecture

# Load Generator



# Load Generator


- Closed loop test ie new job arrivals are only triggered by job completions.
- N threads created to simulate N concurrent users. Each of the N threads emulate one user by issuing one request, waiting for it to complete, and issues the next request immediately afterwards. There is no think time between requests.
- Load generator can generate two types of load (or requests):  
**To create new user account** and **To request file from server1**.
- For each N, we defined either total runtime ie 100-120 seconds or defined number of total fetch file requests. For instance, 2000 file request for each user. So,  $N=1$ , there were total 2000 requests, for  $N=2$ , total 4000 requests, for  $N=10$ , total 20000 requests and so on.

## Phase2 bottleneck


Load type	File size	N*	Mean Response Time at N=N*	Max Throughput at N=N*	Bottleneck
new		800	0.219622	3363.49	Disk I/O (Write) and MySQL
get	1KB	17	0.0054039	3093.7	Disk I/O (Read)
get	10KB	10	0.00390524	2547.3	Disk I/O (Read)
get	100KB	5	0.0113418	430.355	Disk I/O (Read)
get	1MB	2	0.0265555	75.1187	Disk I/O (Read)
get	10MB	1	0.0984429	10.1575	Disk I/O (Read)
get	100MB	1	0.880579	1.1356	Disk I/O (Read)

Table: Test results before optimisation

# Phase2 bottleneck



100KB



1MB

# Phase2 bottleneck

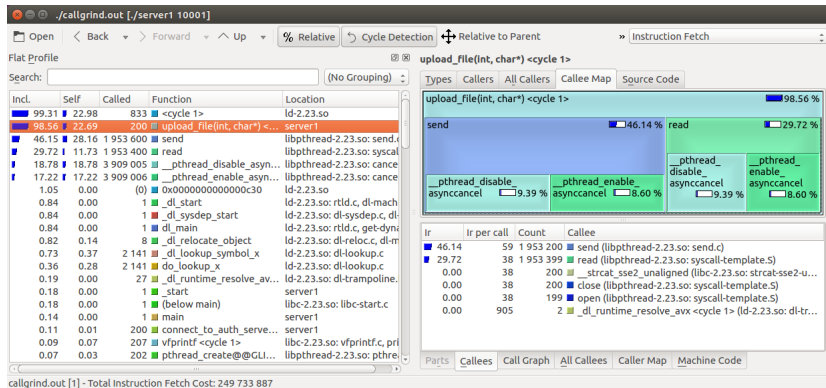


Figure: Profiling the server code



# Optimisation

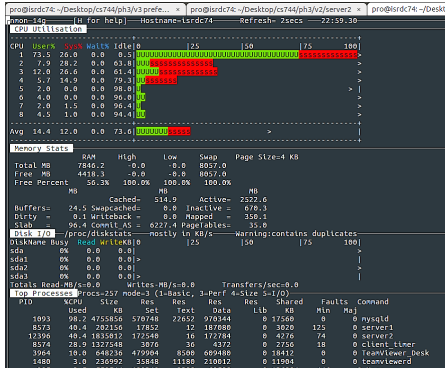
- **Design change.** Changed Server2 from Multi-process to multi-threaded.
- **Disk IO optimization.** Opened files in O\_DIRECT mode and read every file in one I/O, used malloc to store file contents when server started and client request were fetched without disk access.
- **Code optimization.** Keeping variables that will be accessed together (or right after another) next to each other in memory by declaring them together. Also, created structure to include items accessed together. Removing loops to access contents of malloc(ed) files by using memcpy with offset.
- **Database optimisation using Redis Database.** Created username-password as key-value in Redis database and changed server2 code to connect to Redis-server.

# Bottleneck after prefetching using MySQL

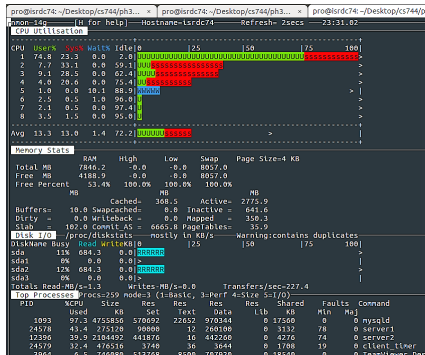
Load type	File size	Max Throughput	Bottleneck
new		6806.03	Disk I/O (Write) and MySQL
get	1KB	9932.58	MySQL
get	10KB	9973.76	MySQL
get	100KB	9804.96	MySQL/ CPU
get	1MB	1633.88	CPU
get	10MB	176.932	CPU
get	100MB	18.0733	CPU

**Table:** Test results after prefetching using MySQL

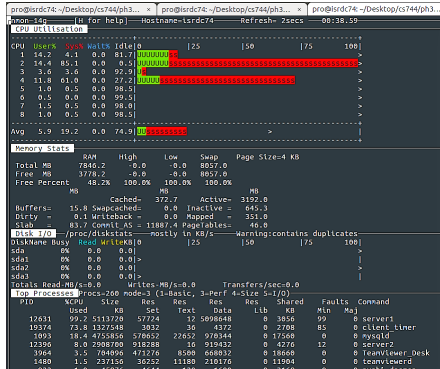
# Bottleneck after prefetching using MySQL



1KB



10KB



100KB

1MB

# Bottleneck after prefetching using MySQL

1KB

```
Totals Read-MB/s=0.0      Writes-MB/s=0.0      Transfers/sec=0.0
Top Processes Procs=257 mode=3 (1-Basic, 3=Perf 4=Size 5=I/O)
```

PID	KCPU	Size	Res	Set	Text	Data	Lib	Shared	Faults	Command
	Used	KB	Res	Text	Data	Lib	KB	Min	Maj	
1093	98.2	4755856	570748	22652	970344	0	17560	0	0	mysqld
8573	40.4	202156	17852	12	187000	0	3070	125	0	server1
12396	40.4	1835012	172540	16	172784	0	4276	74	0	server2
8574	28.9	1327548	3076	36	4372	0	2756	18	0	client_timer
3964	10.0	648236	479904	8500	609480	0	18412	0	0	TeamViewer_Desk
1480	3.0	236992	35848	11180	210012	0	11904	0	0	teamviewerd

10KB

```
Totals Read-MB/s=1.3      Writes-MB/s=0.0      Transfers/sec=227.4
Top Processes Procs=259 mode=3 (1-Basic, 3=Perf 4=Size 5=I/O)
```

PID	KCPU	Size	Res	Set	Text	Data	Lib	Shared	Faults	Command
	Used	KB	Res	Text	Data	Lib	KB	Min	Maj	
1093	97.3	4755856	570692	22652	970344	0	17560	0	0	mysqld
24578	43.4	275120	90000	12	260100	0	3132	78	0	server1
12396	39.9	2104492	441876	16	442260	0	4276	74	0	server2
24579	32.4	476516	3740	36	3644	0	1708	19	0	client_timer
3964	6.5	746080	513768	8500	707920	0	18560	0	0	TeamViewer_Desk

100KB

```
Totals Read-MB/s=0.0      Writes-MB/s=0.0      Transfers/sec=0.0
Top Processes Procs=261 mode=3 (1-Basic, 3=Perf 4=Size 5=I/O)
```

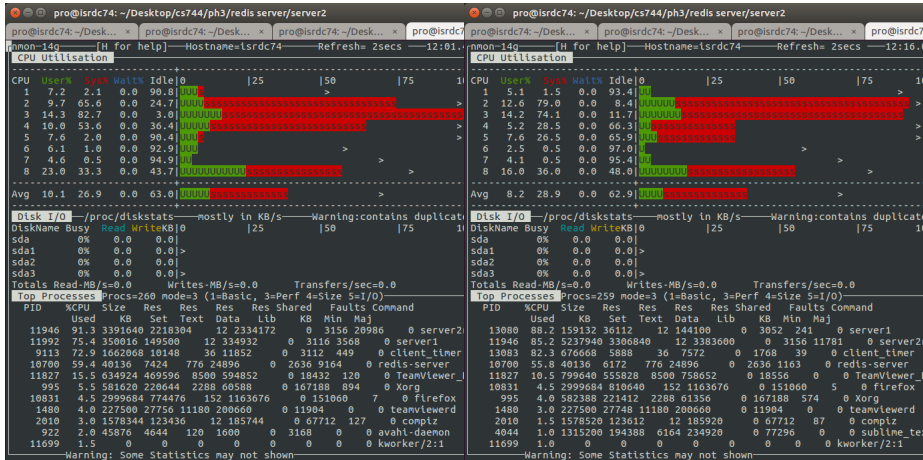
PID	KCPU	Size	Res	Set	Text	Data	Lib	Shared	Faults	Command
	Used	KB	Res	Text	Data	Lib	KB	Min	Maj	
1093	95.7	4755856	570652	22652	970344	0	17560	0	0	mysqld
29008	91.7	1247816	192436	12	1232772	0	3084	836	0	server1
7490	69.3	1003428	5888	36	7292	0	1708	19	0	client_timer
12396	37.9	2398192	735444	16	735960	0	4276	75	0	server2
3964	3.5	716384	482484	8500	680200	0	18660	0	0	TeamViewer_Desk

1MB

```
Totals Read-MB/s=0.0      Writes-MB/s=0.0      Transfers/sec=0.0
Top Processes Procs=260 mode=3 (1-Basic, 3=Perf 4=Size 5=I/O)
```

PID	KCPU	Size	Res	Set	Text	Data	Lib	Shared	Faults	Command
	Used	KB	Res	Text	Data	Lib	KB	Min	Maj	
12631	99.2	5113720	57724	12	5098648	0	3056	99	0	server1
19374	73.8	1327548	3032	36	4372	0	2708	85	0	client_timer
1093	18.4	4755856	570652	22652	970344	0	17560	0	0	mysqld
12396	8.0	2908700	918288	16	919432	0	4276	12	0	server2
3964	3.5	704096	471276	8500	668032	0	18660	0	0	TeamViewer_Desk
1480	1.5	237156	36252	11180	210176	0	11904	0	0	teamviewerd

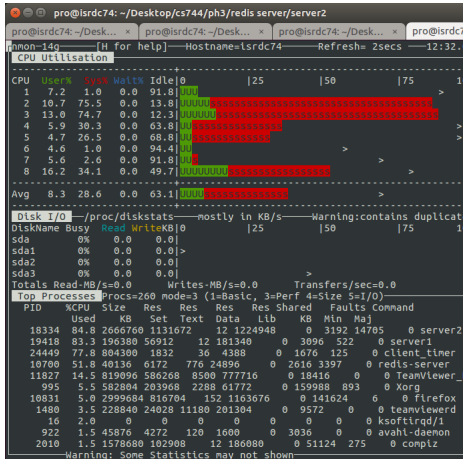
## Bottleneck after prefetching using Redis




1KB

10KB

# Bottleneck after prefetching using Redis



100KB



1MB





## Summary of results

Load type	File size	Max Throughput			
		MT-MP, no prefetching, MySQL	MT-MT, no prefetching, MySQL	MT-MT, prefetching, MySQL	MT-MT, prefetching, Redis
new		3363.49	6806.03	6806.03	23089.8
get	1KB	1136.3921	2556.86	9932.58	23154.2
get	10KB	914.86	1559.36	9973.76	20907.1
get	100KB	697.947	832.57	9804.96	9604.82
get	1MB	132.954	136.1965	1633.88	1624.99
get	10MB	16.8549	16.9793	176.932	175.399
get	100MB	1.9439	1.95548	18.0733	17.6434




Table: Summary of results



# Summary of results



# Summary of results



# Lessons learnt

- Multi-threaded systems give better performance when there is no shared data and overhead of synchronization.
- Pre-fetching the files into user space by allocating memory to file contents using malloc removes disk as bottleneck and increases throughput.
- In-memory database like Redis server gives better throughput performance for small file size when MySQL acts as bottleneck.
- Difficult to remove network bottleneck as it needs hardware changes.
- For optimising performance and identifying hardware bottlenecks, there is requirement to tune sockets, system and MySQL parameters.

# Thank You!